

## sutton-barto-reinforcement-exercises-ch-2

August 26, 2018

2.1:

When using an  $\epsilon$ -greedy action selection criterion, in the case of two actions and  $\epsilon = 0.5$ , actions would be selected at random half of the time. Of the times when actions are selected at random, the greedy action will be selected half of the time. Meaning the greedy action will be selected a total of 75% of the time.

2.2:

Given that all of the reward estimates are initialized to 0, the first action will be selected arbitrarily, possible at random.

The following actions must have been selected at random as a result of the  $\epsilon$  case:  $A_2, A_5$ .

This can be demonstrated by observing that  $Q_1(1) = 1 > Q_1(2) = 0$  and  $Q_4(2) = \frac{R_2 + R_3 + R_4}{3} = \frac{5}{3} > Q_4(3) = 0$ .

Regarding the question of when random  $\epsilon$  selection could have occurred, it could have occurred at any time step after the first.

2.3:

In the long run  $\epsilon = 0.01$  will achieve the highest cumulative reward and probability of selecting the correct answer. The disadvantage of  $\epsilon = 0.01$  relative to  $\epsilon = 0.1$  is that it will cause our agent to take longer to find the optimal action, but it will select that action more consistently once it has been found.

The optimal action has a reward of 1.55. The expected reward per action once the agent has found the best possible action is equal to the sum of the greedy case and the  $\epsilon$  random case:  $1.55(1 - \epsilon) + \epsilon \sum_{i=1}^{10} \frac{q_*(i)}{10}$ .

So the difference between the expected long term reward per action in the  $\epsilon = .01$  and  $\epsilon = .1$  cases is  $1.55(.1 - .01) = 0.1395$ .

2.4:

If  $\alpha$  is not constant, but instead varies with each update step  $n$ , then our incremental update rule utilizing  $\alpha_n$  can be expressed as follows:

$$Q_{n+1} = Q_n + \alpha_n[R_n - Q_n] \quad (1)$$

$$= \alpha_n R_n + (1 - \alpha_n) Q_n \quad (2)$$

$$= \alpha_n R_n + (1 - \alpha_n)(\alpha_{n-1} R_{n-1} + (1 - \alpha_{n-1}) Q_{n-1}) \quad (3)$$

$$= \alpha_n R_n + (1 - \alpha_n) \alpha_{n-1} R_{n-1} + (1 - \alpha_{n-1})(1 - \alpha_n)(Q_{n-2} + \alpha_{n-2}[R_{n-2} - Q_{n-2}]) \quad (4)$$

$$= \alpha_n R_n + (1 - \alpha_n) \alpha_{n-1} R_{n-1} + (1 - \alpha_{n-1})(1 - \alpha_n)((1 - \alpha_{n-2}) Q_{n-2} + \alpha_{n-2} R_{n-2}) \quad (5)$$

$$= \alpha_n R_n + (1 - \alpha_n) \alpha_{n-1} R_{n-1} + (1 - \alpha_{n-1})(1 - \alpha_n) \alpha_{n-2} R_{n-2} + \quad (6)$$

$$\dots + Q_1 \prod_{i=1}^n (1 - \alpha_i) + \alpha_1 R_1 \prod_{j=1}^{n-1} (1 - \alpha_{n-j+1}) \quad (7)$$

$$= Q_1 \prod_{i=1}^n (1 - \alpha_i) + \alpha_n R_n + \sum_{i=1}^{n-1} \alpha_i R_i \prod_{j=1}^{n-i} (1 - \alpha_{n-j+1}) \quad (8)$$

2.5:

```
In [146]: import numpy as np
import random
import numpy.random
import matplotlib.pyplot as plt

# EXERCISE 2.5

steps = 10000
runs = 2000

"""
returns data in the format

{
    "rewards": np.array([1.5, .2, ...]),
    "optimal": np.array([1, 0, ...])
}
"""
def ten_armed_bandit(sample_average = False):
    q = np.zeros(10)
    Q = np.zeros(10)

    rewards = []
    optimal = []

    for i in range(steps):
        arg_max_reward = None
        arg_max_actions = []

        for j in range(len(Q)):
```

```

        estimated_reward = Q[j]

        if arg_max_reward == None or estimated_reward > arg_max_reward:
            arg_max_reward = estimated_reward
            arg_max_actions = [j]

        elif estimated_reward == arg_max_reward:
            arg_max_actions.append(j)

    if random.uniform(0, 1) < .1:
        selected_action = random.choice(range(10))
    else:
        selected_action = random.choice(arg_max_actions)

    actual_reward = q[selected_action]
    alpha = 1/(i + 1) if sample_average else .1
    Q[selected_action] = Q[selected_action] + alpha * (actual_reward -
Q[selected_action])

    was_action_optimal = 1 if actual_reward == max(q) else 0

    rewards.append(actual_reward)
    optimal.append(was_action_optimal)

    random_walk_values = np.random.normal(0, .01, len(q))

    for j in range(len(q)):
        q[j] = q[j] + random_walk_values[j]

    return {
        "rewards": np.array(rewards),
        "optimal": np.array(optimal)
    }

def average_bandit_results(sample_average = False):
    rewards = np.zeros(steps)
    optimal = np.zeros(steps)

    for i in range(runs):
        results = ten_armed_bandit(sample_average)
        rewards = rewards + (results["rewards"] / runs)
        optimal = optimal + (results["optimal"] / runs)

    return {
        "rewards": rewards,
        "optimal": optimal
    }

def plot_average_rewards(with_sample_average, without_sample_average):
    plt.plot(with_sample_average, "r--", label = "sample average")
    plt.plot(without_sample_average, "g", label = "constant")
    plt.xlabel("steps")
    plt.ylabel("average reward")
    plt.legend()
    plt.show()

def plot_average_optimality(with_sample_average, without_sample_average):
    plt.plot(with_sample_average, "r--", label = "sample average")
    plt.plot(without_sample_average, "g", label = "constant")
    plt.xlabel("steps")
    plt.ylabel("percent optimal")
    plt.legend()
    plt.show()


```

In [78]: %matplotlib inline


In [147]: with\_sample\_average = average\_bandit\_results(True)

```
In [148]: without_sample_average = average_bandit_results()
```

```
In [149]: plot_average_rewards(with_sample_average["rewards"], without_sample_average["rewards"])
```



```
In [150]: plot_average_optimality(with_sample_average["optimal"],  
                                   without_sample_average["optimal"])
```



2.6:

The optimistic greedy agent's initial spike in rewards can be explained as follows.

$q_*(a)$  is normally distributed with a mean of 0 and a variance of 1. This means  $Q_1 > \max\{q_*(a)\}$  in the vast majority of cases.

For the reward distributions where the above is true, agents with the optimistic greedy action-value method will attempt every action at least once at the onset.

The spike we observe is a product of agents selecting the optimal action once during this preliminary period.

2.7